

BRAIN TUMOR ANOMALY DETECTION USING AUTOENCODERS

Libraries

```
import kagglehub
import os
import shutil
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras import layers, models
from sklearn.model_selection import train_test_split
from tqdm.keras import TqdmCallback
from sklearn.metrics import accuracy_score, confusion_matrix
from tqdm import tqdm
import seaborn as sns
import pandas as pd
```

Loading dataset

```
#download dataset
path = kagglehub.dataset_download("jakeshbohaju/brain-tumor")
print("Path to dataset files:", path)

#move dataset to content folder
colab_content_dir = "/content/brain-tumor-dataset"

if not os.path.exists(colab_content_dir):
    os.makedirs(colab_content_dir)

for item in os.listdir(path):
    source = os.path.join(path, item)
    destination = os.path.join(colab_content_dir, item)
    if os.path.isdir(source):
        if os.path.exists(destination):
            shutil.rmtree(destination)
        shutil.copytree(source, destination)
    else:
        shutil.copy2(source, destination)
print(f"Moved dataset files to: {colab_content_dir}")
```

Downloading from
https://www.kaggle.com/api/v1/datasets/download/jakeshbohaju/brain-tumor?dataset_version_number=3...

100%|██████████| 14.0M/14.0M [00:01<00:00, 7.68MB/s]

Extracting files...

Path to dataset files:

/root/.cache/kagglehub/datasets/jakeshbohaju/brain-tumor/versions/3

Moved dataset files to: /content/brain-tumor-dataset

```
import pandas as pd
```

```
df = pd.read_csv("/content/brain-tumor-dataset/Brain Tumor.csv")
```

```
df.head()
```

```
{
  "summary": {
    "name": "df",
    "rows": 3762,
    "fields": [
      {
        "column": "Image",
        "properties": {
          "dtype": "string",
          "num_unique_values": 3762,
          "samples": [
            "Image1554",
            "Image2988",
            "Image221"
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "Class",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 0,
          "max": 1,
          "num_unique_values": 2,
          "samples": [
            1,
            0
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "Mean",
        "properties": {
          "dtype": "number",
          "std": 5.728021863519244,
          "min": 0.078659057617188,
          "max": 33.2399749755859,
          "num_unique_values": 3692,
          "samples": [
            8.89088439941406,
            22.7763671875
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "Variance",
        "properties": {
          "dtype": "number",
          "std": 467.46689568736934,
          "min": 3.14562750498484,
          "max": 2910.58187868159,
          "num_unique_values": 3699,
          "samples": [
            206.239411639648,
            612.628337511462
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "Standard Deviation",
        "properties": {
          "dtype": "number",
          "std": 1.77359169624377,
          "min": 0.07026869268617432,
          "max": 53.9498088845697,
          "num_unique_values": 3699,
          "samples": [
            14.3610379722236,
            24.7513300150004
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "Entropy",
        "properties": {
          "dtype": "number",
          "std": 0.07026869268617432,
          "min": 0.000881579569793,
          "max": 0.394538600726663,
          "num_unique_values": 3699,
          "samples": [
            0.102807935605076,
            0.070073067535855
          ],
          "semantic_type": "",
          "description": ""
        }
      ]
    }
  }
}
```

```

\"Skewness\",\\n      \"properties\": {\\n      \"dtype\":
\"number\",\\n      \"std\": 2.5609398180749827,\\n      \"min\":
1.88601442072597,\\n      \"max\": 36.9312940533355,\\n
\"num_unique_values\": 3699,\\n      \"samples\": [\\n
3.32352991443732,\\n      2.75330347584502\\n      ],\\n
\"semantic_type\": \"\",\\n      \"description\": \"\",\\n      }\\n
},\\n      {\\n      \"column\": \"Kurtosis\",\\n      \"properties\":
{\\n      \"dtype\": \"number\",\\n      \"std\":
56.434747010233224,\\n      \"min\": 3.94240208400556,\\n
\"max\": 1371.6400603465,\\n      \"num_unique_values\": 3699,\\n
\"samples\": [\\n      11.335489542349,\\n
7.84468707653204\\n      ],\\n      \"semantic_type\": \"\",\\n
\"description\": \"\",\\n      }\\n      },\\n      {\\n      \"column\":
\"Contrast\",\\n      \"properties\": {\\n      \"dtype\":
\"number\",\\n      \"std\": 109.49960055870405,\\n      \"min\":
3.19473319473319,\\n      \"max\": 3382.57416267943,\\n
\"num_unique_values\": 3699,\\n      \"samples\": [\\n
68.1793591344153,\\n      120.88492808005\\n      ],\\n
\"semantic_type\": \"\",\\n      \"description\": \"\",\\n      }\\n
},\\n      {\\n      \"column\": \"Energy\",\\n      \"properties\":
{\\n      \"dtype\": \"number\",\\n      \"std\":
0.12935164441107974,\\n      \"min\": 0.024731170883341,\\n
\"max\": 0.589681787363579,\\n      \"num_unique_values\": 3699,\\n
\"samples\": [\\n      0.283209438335796,\\n
0.231722468697145\\n      ],\\n      \"semantic_type\": \"\",\\n
\"description\": \"\",\\n      }\\n      },\\n      {\\n      \"column\":
\"ASM\",\\n      \"properties\": {\\n      \"dtype\": \"number\",\\n
\"std\": 0.05830035624005218,\\n      \"min\": 0.000611630813261,\\n
\"max\": 0.347724610348305,\\n      \"num_unique_values\": 3699,\\n
\"samples\": [\\n      0.080207585962477,\\n
0.0536953024991\\n      ],\\n      \"semantic_type\": \"\",\\n
\"description\": \"\",\\n      }\\n      },\\n      {\\n      \"column\":
\"Homogeneity\",\\n      \"properties\": {\\n      \"dtype\":
\"number\",\\n      \"std\": 0.12792908048929705,\\n      \"min\":
0.105489790279749,\\n      \"max\": 0.810920845803123,\\n
\"num_unique_values\": 3699,\\n      \"samples\": [\\n
0.585790793249951,\\n      0.511982977953878\\n      ],\\n
\"semantic_type\": \"\",\\n      \"description\": \"\",\\n      }\\n
},\\n      {\\n      \"column\": \"Dissimilarity\",\\n      \"properties\":
{\\n      \"dtype\": \"number\",\\n      \"std\":
1.8501726110682728,\\n      \"min\": 0.681120681120681,\\n
\"max\": 27.8277511961722,\\n      \"num_unique_values\": 3698,\\n
\"samples\": [\\n      3.30898876404494,\\n
4.70559724828018\\n      ],\\n      \"semantic_type\": \"\",\\n
\"description\": \"\",\\n      }\\n      },\\n      {\\n      \"column\":
\"Correlation\",\\n      \"properties\": {\\n      \"dtype\":
\"number\",\\n      \"std\": 0.026156805657855237,\\n      \"min\":
0.549426249103514,\\n      \"max\": 0.989972351110618,\\n
\"num_unique_values\": 3699,\\n      \"samples\": [\\n

```

```
0.941727270038754,\n          0.948646654751285\n          ],\n  \"semantic_type\": \"\",\n  \"description\": \"\"\n  }\n  },\n  {\n    \"column\": \"Coarseness\",\n    \"properties\": {\n      \"dtype\": \"number\",\n      \"std\":\n      0.0,\n      \"min\": 7.45834073119875e-155,\n      \"max\":\n      7.45834073120021e-155,\n      \"num_unique_values\": 146,\n      \"samples\": [\n        7.458340731199201e-155,\n        7.458340731199731e-155\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\"\n    }\n  }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"df\"}
```

```
df['Class'].value_counts()
```

```
Class
```

```
0    2079
```

```
1    1683
```

```
Name: count, dtype: int64
```

Class distribution

```
# Plot with seaborn
plt.figure(figsize=(8, 5))
sns.countplot(x='Class', data=df, palette='Set2', alpha=0.7) # 'Set2'
provides multiple colors
plt.title('Class Distribution of Benign and Malignant Tumor')
plt.xlabel('Class')
plt.ylabel('Frequency')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```


```
print("Class Distribution:")
```

```
print(df['Class'].value_counts())
```

```
<ipython-input-28-7838effe8039>:3: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Class', data=df, palette='Set2', alpha=0.7) #
'Set2' provides multiple colors
```



```
Class Distribution:
Class
0    2079
1    1683
Name: count, dtype: int64
```

Loading samples

```
#images folder
image_folder = "/content/brain-tumor-dataset/Brain Tumor/Brain Tumor/"

#load image function
def load_images_with_features(folder_path, dataframe,
num_images_to_show=15):
    images = []
    features = []
    for i, filename in enumerate(os.listdir(folder_path)):
        if i >= num_images_to_show:
            break
        img_path = os.path.join(folder_path, filename)
        if os.path.isfile(img_path) and filename.endswith('.jpg'):
            # Load image
            img = Image.open(img_path).convert('L')
            img_array = np.array(img)
```

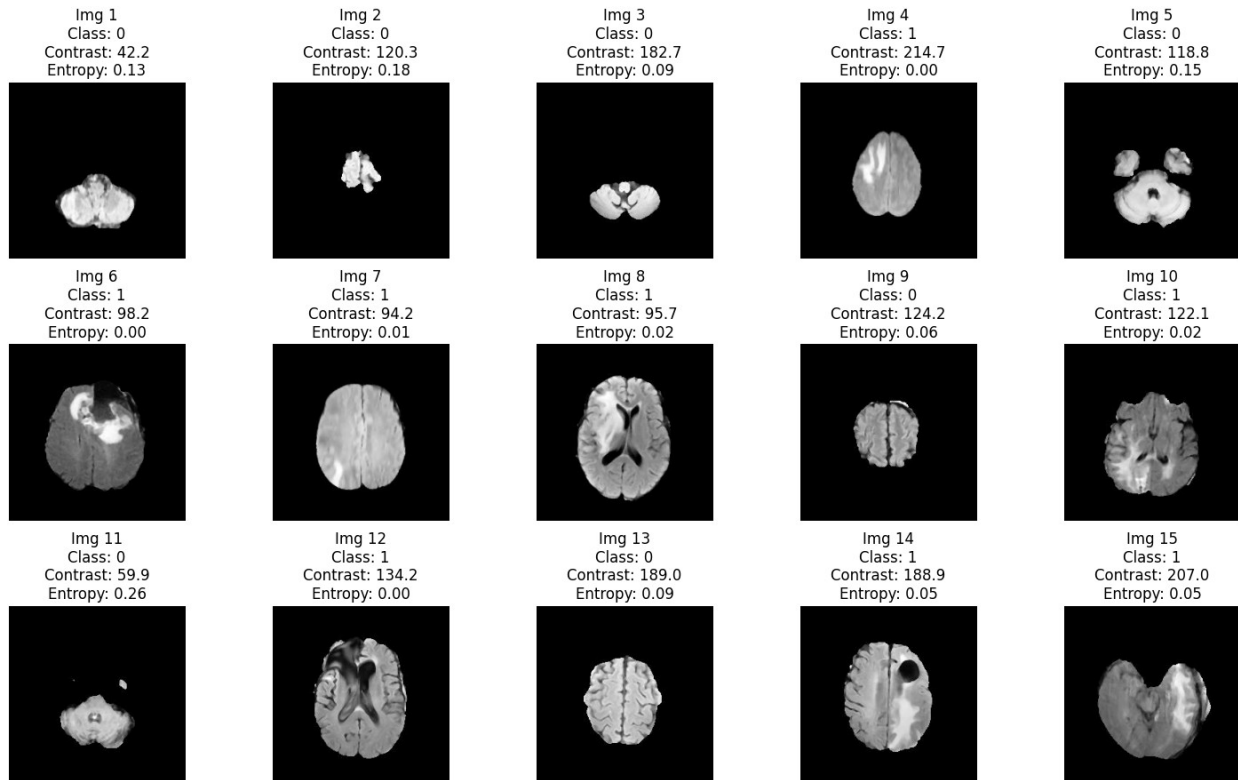
```

        images.append(img_array)
        # Match with CSV (assuming 'Image' column has filenames)
        img_name = filename.split('.')[0] # Remove .jpg extension
        row = dataframe[dataframe['Image'] == img_name]
        if not row.empty:
            features.append(row.iloc[0])
        else:
            features.append(None) # If no match found
    return images, features

#load 15 images
num_images_to_show = 15
images, features = load_images_with_features(image_folder, df,
num_images_to_show)

# plot with attributes
if images:
    plt.figure(figsize=(15, 9))
    for i in range(min(num_images_to_show, len(images))):
        row = i // 5
        col = i % 5
        plt.subplot(3, 5, i + 1)
        plt.imshow(images[i], cmap='gray')
        #add title with Class
        if features[i] is not None:
            title = (f"Img {i+1}\nClass: {features[i]['Class']}\n"
                    f"Contrast: {features[i]['Contrast']:.1f}\n"
                    f"Entropy: {features[i]['Entropy']:.2f}")
        else:
            title = f"Img {i+1}\n(No data)"
        plt.title(title)
        plt.axis('off')
    plt.tight_layout()
    plt.show()
else:
    print("No images loaded.")

```



Splitting data

```
# load images
def load_images_from_folder(folder_path, image_names):
    images = []
    for filename in image_names:
        img_path = os.path.join(folder_path, filename + '.jpg')
        if os.path.isfile(img_path):
            img = Image.open(img_path).convert('L')
            img = img.resize((128, 128))
            img_array = np.array(img) / 255.0
            images.append(img_array)
    return np.array(images)

# Separate Class 0 and Class 1 based on CSV
class_0_df = df[df['Class'] == 0]
class_1_df = df[df['Class'] == 1]

# Train set: 1500 Class 0
train_df = class_0_df.sample(n=1500, random_state=42) # Randomly
select 1500 Class 0
train_images = load_images_from_folder(image_folder,
train_df['Image'])
train_labels = np.zeros(len(train_images)) # All Class 0

# Test set: 579 Class 0 (remaining) + 570 Class 1
```

```

remaining_class_0_df = class_0_df.drop(train_df.index) # 579
remaining_class_0
test_class_0_df = remaining_class_0_df # All 579
test_class_1_df = class_1_df.sample(n=570, random_state=42) # 570
Class 1

# Load test images
test_class_0_images = load_images_from_folder(image_folder,
test_class_0_df['Image'])
test_class_1_images = load_images_from_folder(image_folder,
test_class_1_df['Image'])

# Combine test set
test_images = np.concatenate([test_class_0_images,
test_class_1_images])
test_labels = np.concatenate([np.zeros(len(test_class_0_images)),
np.ones(len(test_class_1_images))])

# Reshape for model input
train_images = train_images.reshape(-1, 128, 128, 1)
test_images = test_images.reshape(-1, 128, 128, 1)

# Print shapes and label distribution
print(f"Total images loaded: {len(df)}") # Should be 3762
print(f"Training images (Class 0): {train_images.shape}")
print(f"Testing images (mixed): {test_images.shape}, Labels:
{test_labels.shape}")
print("Test label distribution:",
np.bincount(test_labels.astype(int)))

Total images loaded: 3762
Training images (Class 0): (1500, 128, 128, 1)
Testing images (mixed): (1149, 128, 128, 1), Labels: (1149,)
Test label distribution: [579 570]

```

Autoencoder

```

#gpu
gpus = tf.config.list_physical_devices('GPU')
if gpus:
    print(f"GPUs available: {len(gpus)}")
    for gpu in gpus:
        print(f"GPU: {gpu}")
        tf.config.experimental.set_memory_growth(gpu, True)
else:
    print("No GPU available, running on CPU.")

#Autoencoder

```



```
input_img = layers.Input(shape=(128, 128, 1))

#Encoder
x = layers.Conv2D(32, (3, 3), activation='relu', padding='same')(input_img)
x = layers.MaxPooling2D((2, 2), padding='same')(x)
x = layers.Conv2D(16, (3, 3), activation='relu', padding='same')(x)
encoded = layers.MaxPooling2D((2, 2), padding='same')(x)

#Decoder
x = layers.Conv2D(16, (3, 3), activation='relu', padding='same')(encoded)
x = layers.UpSampling2D((2, 2))(x)
x = layers.Conv2D(32, (3, 3), activation='relu', padding='same')(x)
x = layers.UpSampling2D((2, 2))(x)
decoded = layers.Conv2D(1, (3, 3), activation='sigmoid', padding='same')(x)

#initialization
autoencoder = models.Model(input_img, decoded)

#optimizer
policy = tf.keras.mixed_precision.Policy('mixed_float16')
tf.keras.mixed_precision.set_global_policy(policy)
autoencoder.compile(optimizer='adam', loss='mean_squared_error')

#summary
autoencoder.summary()

#split training set 80% for train and 20% for validation
train_data, val_data = train_test_split(train_images, test_size=0.2, random_state=42)

#create tf.data.Dataset for training
train_dataset = tf.data.Dataset.from_tensor_slices((train_data, train_data))
train_dataset = train_dataset.shuffle(buffer_size=1024).batch(32).prefetch(tf.data.AUTOTUNE)

#create tf.data.Dataset for validation
val_dataset = tf.data.Dataset.from_tensor_slices((val_data, val_data))
val_dataset = val_dataset.batch(32).prefetch(tf.data.AUTOTUNE)

#train autoencoder with class 0
with tf.device('/GPU:0'):
    history = autoencoder.fit(train_dataset, epochs=100, validation_data=val_dataset, callbacks=[TqdmCallback(verbose=1)],
```

verbose=0)

#training loss

```
plt.figure(figsize=(10, 5))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Autoencoder Training Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

#prediction

```
test_predictions = autoencoder.predict(test_images)
```

#reconstruction error per images

```
reconstruction_errors = np.mean((test_images - test_predictions) ** 2,
axis=(1, 2, 3))
```

#separate errors by true class (using test_labels)

```
class_0_errors = reconstruction_errors[test_labels == 0] # 579 Class
0
class_1_errors = reconstruction_errors[test_labels == 1] # 570 Class
1
```

GPUs available: 1

GPU: PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')

Model: "functional"

Layer (type) Param #	Output Shape
input_layer (InputLayer) 0	(None, 128, 128, 1)
conv2d (Conv2D) 320	(None, 128, 128, 32)
max_pooling2d (MaxPooling2D) 0	(None, 64, 64, 32)
conv2d_1 (Conv2D) 4,624	(None, 64, 64, 16)

0	max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 16)
2,320	conv2d_2 (Conv2D)	(None, 32, 32, 16)
0	up_sampling2d (UpSampling2D)	(None, 64, 64, 16)
4,640	conv2d_3 (Conv2D)	(None, 64, 64, 32)
0	up_sampling2d_1 (UpSampling2D)	(None, 128, 128, 32)
289	conv2d_4 (Conv2D)	(None, 128, 128, 1)

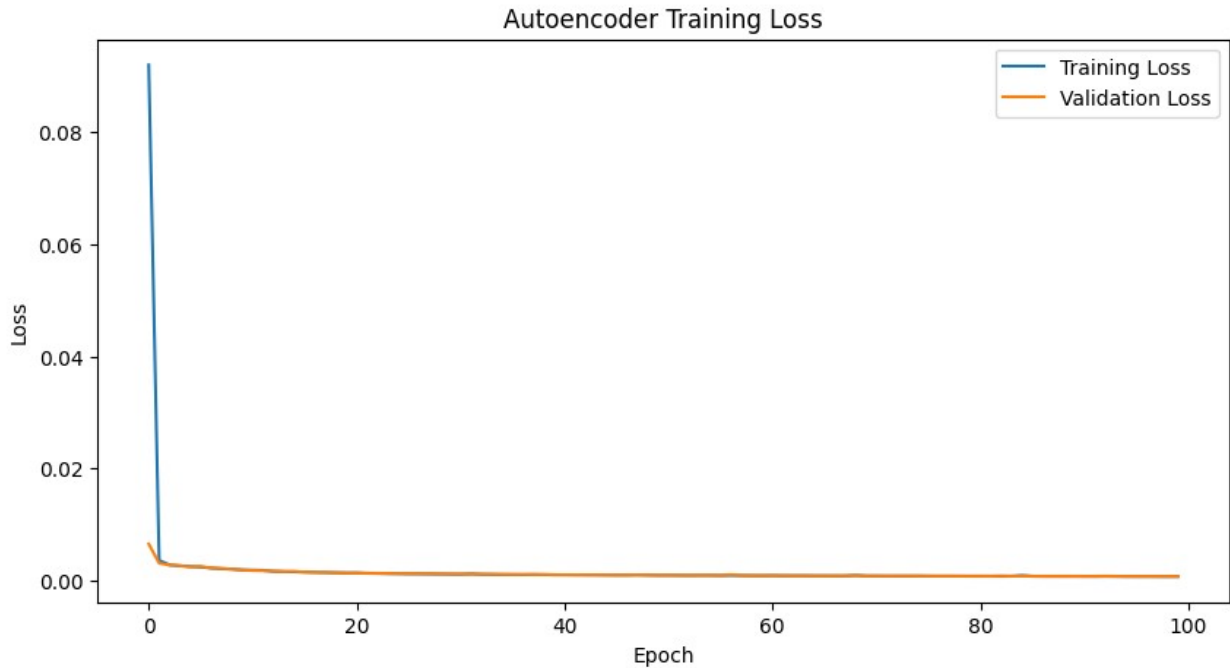
Total params: 12,193 (47.63 KB)

Trainable params: 12,193 (47.63 KB)

Non-trainable params: 0 (0.00 B)

```
{"model_id":"d7acde0af6244a9794fc6c86046e4576","version_major":2,"version_minor":0}
```

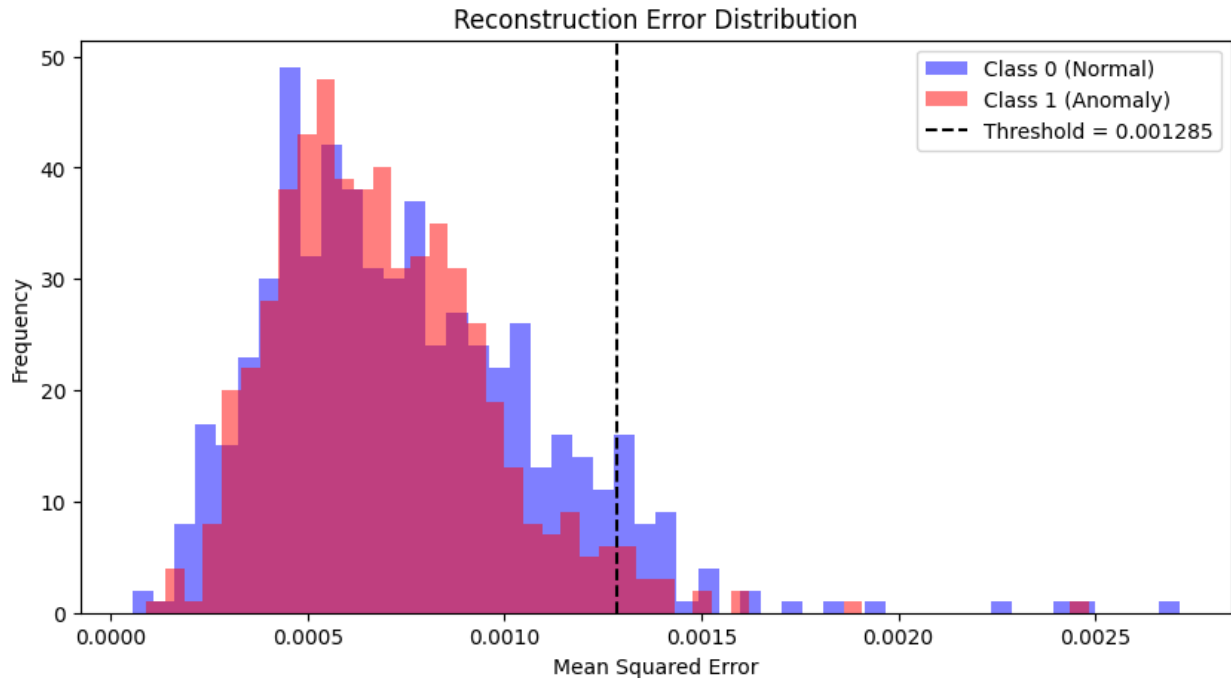
```
{"model_id":"fe6dc334a1f845768feb14ce65c16b46","version_major":2,"version_minor":0}
```



36/36 ————— 2s 39ms/step

Autoencoder reconstruction error distribution

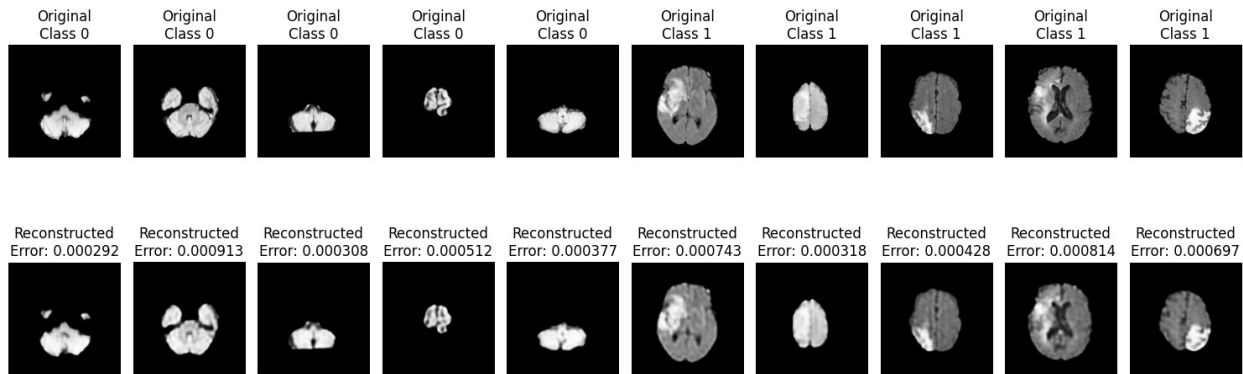
```
#reconstruction errors distribution
threshold = np.mean(class_0_errors) + 1.5 * np.std(class_0_errors)
plt.figure(figsize=(10, 5))
plt.hist(class_0_errors, bins=50, alpha=0.5, label='Class 0 (Normal)',
color='blue')
plt.hist(class_1_errors, bins=50, alpha=0.5, label='Class 1
(Anomaly)', color='red')
plt.axvline(threshold, color='black', linestyle='--',
label=f'Threshold = {threshold:.6f}')
plt.title('Reconstruction Error Distribution')
plt.xlabel('Mean Squared Error')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```



```
#reconstruction error for 5 images of each class
num_examples = 5
class_0_indices = np.where(test_labels == 0)[0][:num_examples]
class_1_indices = np.where(test_labels == 1)[0][:num_examples]

plt.figure(figsize=(15, 6))
for i, idx in enumerate(np.concatenate([class_0_indices,
class_1_indices])):
    plt.subplot(2, num_examples * 2, i + 1)
    plt.imshow(test_images[idx].reshape(128, 128), cmap='gray')
    plt.title(f"Original\nClass {int(test_labels[idx])}")
    plt.axis('off')

    plt.subplot(2, num_examples * 2, i + 1 + num_examples * 2)
    plt.imshow(test_predictions[idx].reshape(128, 128), cmap='gray')
    plt.title(f"Reconstructed\nError:
{reconstruction_errors[idx]:.6f}")
    plt.axis('off')
plt.tight_layout()
plt.show()
```



Adjusting threshold 1-100% for accuracy

```
#reconstruct test images
reconstructed_images = autoencoder.predict(test_images)

#compute reconstruction error (MSE) for each image
mse = np.mean(np.square(test_images - reconstructed_images), axis=(1, 2, 3))

#compute training MSE for threshold calculation
train_reconstructed = autoencoder.predict(train_images)
train_mse = np.mean(np.square(train_images - train_reconstructed), axis=(1, 2, 3))

#dictionary to store threshold and accuracy results
results = {}

#test thresholds from 1% to 100%
for i in tqdm(range(100), desc="Evaluating Thresholds"):
    percentile = 1 + i
    threshold = np.percentile(train_mse, percentile) # Percentile of Class 0 errors

    # Classify test images as normal (0) or anomalous (1) based on threshold
    predictions = (mse > threshold).astype(int)

    # Evaluate performance
    accuracy = accuracy_score(test_labels, predictions)
    conf_matrix = confusion_matrix(test_labels, predictions)

    # Store results
    results[percentile] = {
        'threshold': threshold,
        'accuracy': accuracy,
        'conf_matrix': conf_matrix
    }
```

```

# After all iterations, print all results and find the best one
print("\n=== All Results ===")
best_percentile = None
best_accuracy = -1
for percentile, result in results.items():
    if result['accuracy'] > best_accuracy:
        best_accuracy = result['accuracy']
        best_percentile = percentile

print("\n=== Best Result ===")
print(f"Best Percentile: {best_percentile}%")
print(f"Threshold: {results[best_percentile]['threshold']:.6f}")
print(f"Accuracy: {results[best_percentile]['accuracy']:.4f}")
print(f"Confusion Matrix:\n{results[best_percentile]['conf_matrix']}")

# Plot percentile vs accuracy
percentiles_ae = list(results.keys()) # [1, 2, ..., 100]
accuracies_ae = [results[p]['accuracy'] for p in results]

plt.figure(figsize=(10, 6))
plt.plot(percentiles_ae, accuracies_ae, marker='o', linestyle='--',
color='g')
plt.title('Percentile vs Accuracy')
plt.xlabel('Percentile (%)')
plt.ylabel('Accuracy')
plt.grid(True)
plt.show()

```

```

36/36 _____ 0s 8ms/step
47/47 _____ 0s 8ms/step

```

```

Evaluating Thresholds: 100%|██████████| 100/100 [00:00<00:00,
347.86it/s]

```

```

=== All Results ===

```

```

=== Best Result ===

```

```

Best Percentile: 4%

```

```

Threshold: 0.000273

```

```

Accuracy: 0.5117

```

```

Confusion Matrix:

```

```

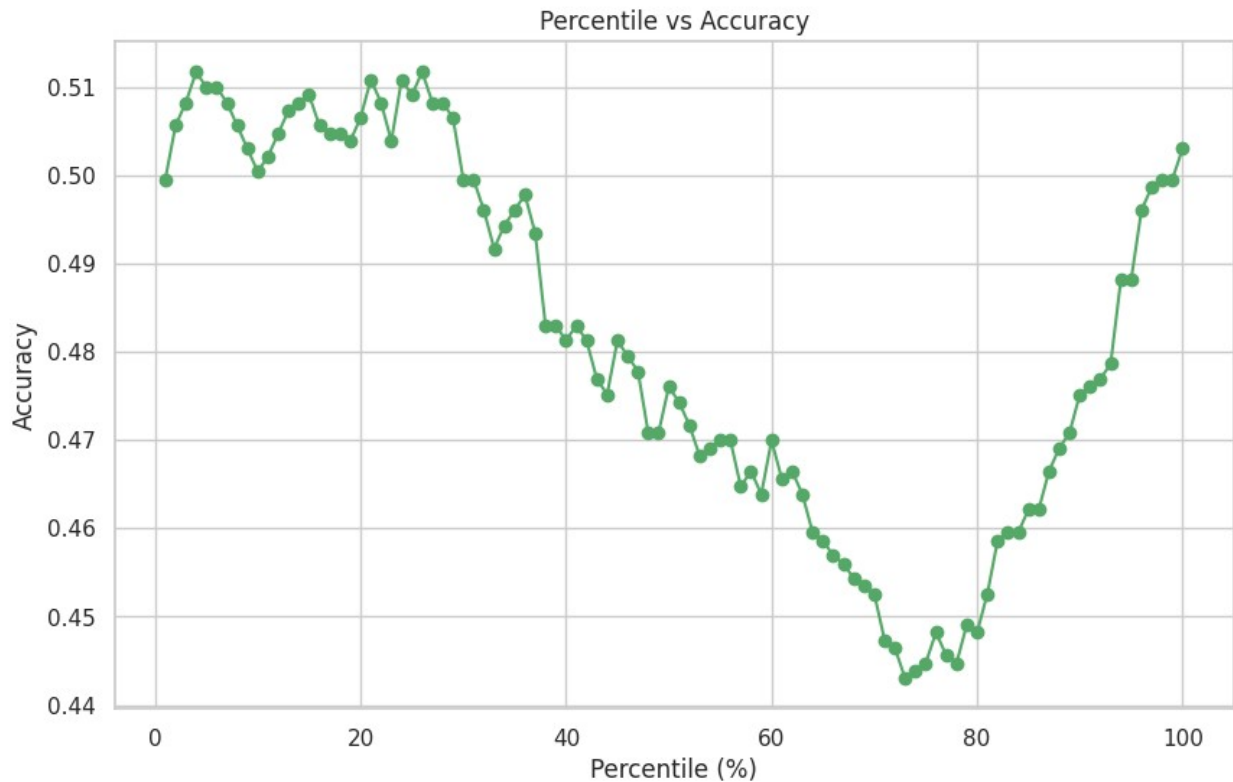
[[ 29 550]

```

```

 [ 11 559]]

```



Best result from Autoencoder

```
import seaborn as sns
# Generate heatmap for the best result's confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(results[best_percentile]['conf_matrix'], annot=True,
            fmt='d', cmap='Blues',
            xticklabels=['Predicted Class 0', 'Predicted Class 1'],
            yticklabels=['True Class 0', 'True Class 1'])
plt.title(f'Confusion Matrix for Best Threshold (Percentile:
{best_percentile}%, Accuracy: {best_accuracy:.4f})')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```


Confusion Matrix for Best Threshold (Percentile: 4%, Accuracy: 0.5117)



Variational Autoencoder

```
# Check GPU availability
gpus = tf.config.list_physical_devices('GPU')
if gpus:
    print(f"GPUs available: {len(gpus)}")
    for gpu in gpus:
        print(f"GPU: {gpu}")
        tf.config.experimental.set_memory_growth(gpu, True)
else:
    print("No GPU available, running on CPU.")

# Custom VAE layer for sampling
class Sampling(layers.Layer):
    def call(self, inputs):
        z_mean, z_log_var = inputs
        batch = tf.shape(z_mean)[0]
```

```

        dim = tf.shape(z_mean)[1]
        epsilon = tf.keras.backend.random_normal(shape=(batch, dim),
dtype=tf.float16)
        return z_mean + tf.cast(tf.exp(tf.cast(0.5 * z_log_var,
tf.float16))), tf.float16) * epsilon

# VAE parameters
latent_dim = 64 # Increased latent space size
input_shape = (128, 128, 1)

# Encoder (deeper)
inputs = layers.Input(shape=input_shape)
x = layers.Conv2D(32, (3, 3), activation='relu', padding='same')(
inputs)
x = layers.MaxPooling2D((2, 2), padding='same')(x)
x = layers.Conv2D(64, (3, 3), activation='relu', padding='same')(x)
x = layers.MaxPooling2D((2, 2), padding='same')(x)
x = layers.Conv2D(128, (3, 3), activation='relu', padding='same')(x)
x = layers.MaxPooling2D((2, 2), padding='same')(x)
x = layers.Conv2D(256, (3, 3), activation='relu', padding='same')(x)
# Added layer
x = layers.MaxPooling2D((2, 2), padding='same')(x)
x = layers.Flatten()(x)
z_mean = layers.Dense(latent_dim, name='z_mean')(x)
z_log_var = layers.Dense(latent_dim, name='z_log_var')(x)
z = Sampling()([z_mean, z_log_var])
encoder = models.Model(inputs, [z_mean, z_log_var, z], name='encoder')

# Decoder (deeper)
latent_inputs = layers.Input(shape=(latent_dim,))
x = layers.Dense(8 * 8 * 256)(latent_inputs) # Adjusted for deeper
encoder
x = layers.Reshape((8, 8, 256))(x)
x = layers.Conv2D(256, (3, 3), activation='relu', padding='same')(x)
x = layers.UpSampling2D((2, 2))(x)
x = layers.Conv2D(128, (3, 3), activation='relu', padding='same')(x)
x = layers.UpSampling2D((2, 2))(x)
x = layers.Conv2D(64, (3, 3), activation='relu', padding='same')(x)
x = layers.UpSampling2D((2, 2))(x)
x = layers.Conv2D(32, (3, 3), activation='relu', padding='same')(x)
x = layers.UpSampling2D((2, 2))(x)
outputs = layers.Conv2D(1, (3, 3), activation='sigmoid',
padding='same')(x)
decoder = models.Model(latent_inputs, outputs, name='decoder')

# VAE model
class VAE(models.Model):
    def __init__(self, encoder, decoder, **kwargs):
        super(VAE, self).__init__(**kwargs)
        self.encoder = encoder

```

```

        self.decoder = decoder
        self.total_loss_tracker =
tf.keras.metrics.Mean(name="total_loss")
        self.reconstruction_loss_tracker =
tf.keras.metrics.Mean(name="reconstruction_loss")
        self.kl_loss_tracker = tf.keras.metrics.Mean(name="kl_loss")

    def call(self, inputs, training=None):
        z_mean, z_log_var, z = self.encoder(inputs)
        reconstruction = self.decoder(z)
        return reconstruction

    @property
    def metrics(self):
        return [self.total_loss_tracker,
self.reconstruction_loss_tracker, self.kl_loss_tracker]

    def train_step(self, data):
        if isinstance(data, tuple):
            x = data[0]
        else:
            x = data
        with tf.GradientTape() as tape:
            z_mean, z_log_var, z = self.encoder(x)
            reconstruction = self.decoder(z)
            reconstruction_loss =
tf.reduce_mean(tf.reduce_sum(tf.keras.losses.binary_crossentropy(x,
reconstruction), axis=[1, 2]))
            kl_loss = -0.5 * (1 + z_log_var - tf.square(z_mean) -
tf.exp(z_log_var))
            kl_loss = tf.reduce_mean(tf.reduce_sum(kl_loss, axis=1))
            total_loss = reconstruction_loss + 0.1 * kl_loss #
Reduced KL weight
            grads = tape.gradient(total_loss, self.trainable_weights)
            self.optimizer.apply_gradients(zip(grads,
self.trainable_weights))
            self.total_loss_tracker.update_state(total_loss)

self.reconstruction_loss_tracker.update_state(reconstruction_loss)
self.kl_loss_tracker.update_state(kl_loss)
        return {
            "loss": self.total_loss_tracker.result(),
            "reconstruction_loss":
self.reconstruction_loss_tracker.result(),
            "kl_loss": self.kl_loss_tracker.result(),
        }

    def test_step(self, data):
        if isinstance(data, tuple):
            x = data[0]

```

```

        else:
            x = data
            z_mean, z_log_var, z = self.encoder(x)
            reconstruction = self.decoder(z)
            reconstruction_loss =
tf.reduce_mean(tf.reduce_sum(tf.keras.losses.binary_crossentropy(x,
reconstruction), axis=[1, 2]))
            kl_loss = -0.5 * (1 + z_log_var - tf.square(z_mean) -
tf.exp(z_log_var))
            kl_loss = tf.reduce_mean(tf.reduce_sum(kl_loss, axis=1))
            total_loss = reconstruction_loss + 0.1 * kl_loss
            self.total_loss_tracker.update_state(total_loss)

self.reconstruction_loss_tracker.update_state(reconstruction_loss)
self.kl_loss_tracker.update_state(kl_loss)
        return {
            "loss": self.total_loss_tracker.result(),
            "reconstruction_loss":
self.reconstruction_loss_tracker.result(),
            "kl_loss": self.kl_loss_tracker.result(),
        }

# Instantiate and compile VAE
vae = VAE(encoder, decoder)
vae.compile(optimizer='adam')

# Use mixed precision
policy = tf.keras.mixed_precision.Policy('mixed_float16')
tf.keras.mixed_precision.set_global_policy(policy)

# Split train_images into training and validation
train_data, val_data = train_test_split(train_images, test_size=0.2,
random_state=42)

# Create datasets
train_dataset = tf.data.Dataset.from_tensor_slices((train_data,
train_data)).shuffle(1024).batch(32).prefetch(tf.data.AUTOTUNE)
val_dataset = tf.data.Dataset.from_tensor_slices((val_data,
val_data)).batch(32).prefetch(tf.data.AUTOTUNE)

# Train the VAE on Class 0
with tf.device('/GPU:0'):
    history = vae.fit(train_dataset,
                      epochs=100, # Increased epochs
                      validation_data=val_dataset,
                      callbacks=[TqdmCallback(verbose=1)],
                      verbose=0)

# Plot training loss
plt.plot(history.history['loss'], label='Training Loss')

```

```

plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('VAE Training Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()

# Predict reconstructions for test set
_, _, z = encoder.predict(test_images)
test_predictions = decoder.predict(z)

# Compute reconstruction errors
reconstruction_errors = np.mean((test_images - test_predictions) ** 2,
axis=(1, 2, 3))

# Separate errors by class
class_0_errors = reconstruction_errors[test_labels == 0] # 579
class_1_errors = reconstruction_errors[test_labels == 1] # 570

# Set threshold (tighter)
threshold = np.mean(class_0_errors) + 1.5 * np.std(class_0_errors) #
Adjusted to 1.5*std
print(f"Reconstruction error threshold: {threshold:.6f}")

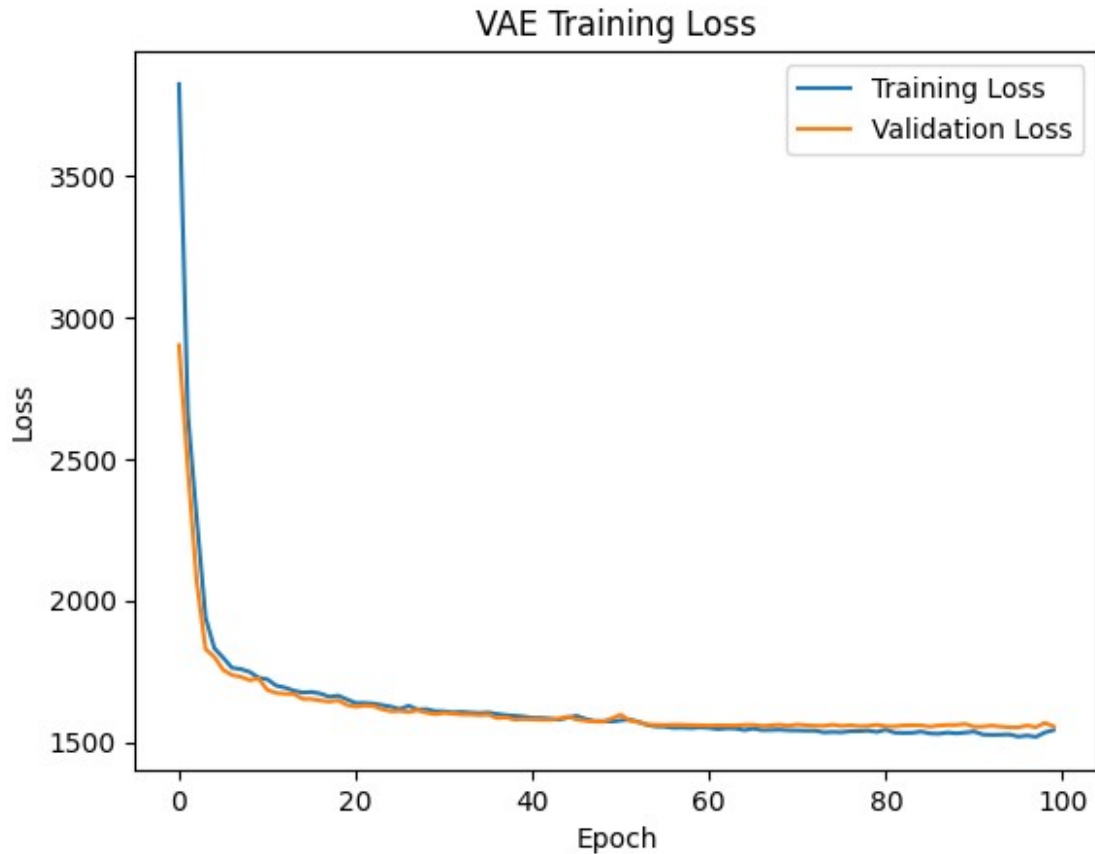
# Classify
predicted_labels = (reconstruction_errors > threshold).astype(int)

GPUs available: 1
GPU: PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')

{"model_id": "70d16836e8aa410a94f7e02bfef73bca", "version_major": 2, "version_minor": 0}

{"model_id": "21d1cb672d9846cd8ba3343fe89dd713", "version_major": 2, "version_minor": 0}

```



36/36 ————— 4s 51ms/step
 36/36 ————— 2s 33ms/step
 Reconstruction error threshold: 0.005066

VAE reconstruction error distribution

```
# Plot error distribution for insight
plt.figure(figsize=(10, 5))
plt.hist(class_0_errors, bins=50, alpha=0.5, label='Class 0',
color='blue')
plt.hist(class_1_errors, bins=50, alpha=0.5, label='Class 1',
color='red')
plt.axvline(threshold, color='black', linestyle='--',
label=f'Threshold = {threshold:.6f}')
plt.title('Reconstruction Error Distribution')
plt.xlabel('Mean Squared Error')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```



```
# Reconstruct test images using VAE's encoder and decoder
_, _, z = encoder.predict(test_images) # Get latent representation
reconstructed_images = decoder.predict(z) # Decode to reconstructions

# Compute reconstruction error (MSE) for each image
mse = np.mean(np.square(test_images - reconstructed_images), axis=(1,
2, 3))

# Set a threshold for anomaly detection (e.g., 95th percentile of
training error)
_, _, train_z = encoder.predict(train_images) # Encode training data
train_reconstructed = decoder.predict(train_z) # Decode training
reconstructions
train_mse = np.mean(np.square(train_images - train_reconstructed),
axis=(1, 2, 3))
threshold = np.percentile(train_mse, 95) # 95th percentile of Class 0
errors
print(f"Anomaly detection threshold: {threshold:.6f}")

# Classify test images as normal (0) or anomalous (1) based on
threshold
predictions = (mse > threshold).astype(int)

# Evaluate performance
accuracy = accuracy_score(test_labels, predictions)
conf_matrix = confusion_matrix(test_labels, predictions)
```

```

print(f"Accuracy: {accuracy:.4f}")
print("Confusion Matrix:")
print(conf_matrix)

# Visualize example reconstructions (first 5 from each class)
num_examples = 5
class_0_indices = np.where(test_labels == 0)[0][:num_examples]
class_1_indices = np.where(test_labels == 1)[0][:num_examples]

plt.figure(figsize=(15, 6))
for i, idx in enumerate(np.concatenate([class_0_indices,
class_1_indices])):
    plt.subplot(2, num_examples * 2, i + 1)
    plt.imshow(test_images[idx].reshape(128, 128), cmap='gray')
    plt.title(f"Original\nClass {int(test_labels[idx])}")
    plt.axis('off')

    plt.subplot(2, num_examples * 2, i + 1 + num_examples * 2)
    plt.imshow(reconstructed_images[idx].reshape(128, 128),
cmap='gray')
    plt.title(f"Reconstructed\nError: {mse[idx]:.6f}")
    plt.axis('off')
plt.tight_layout()
plt.show()

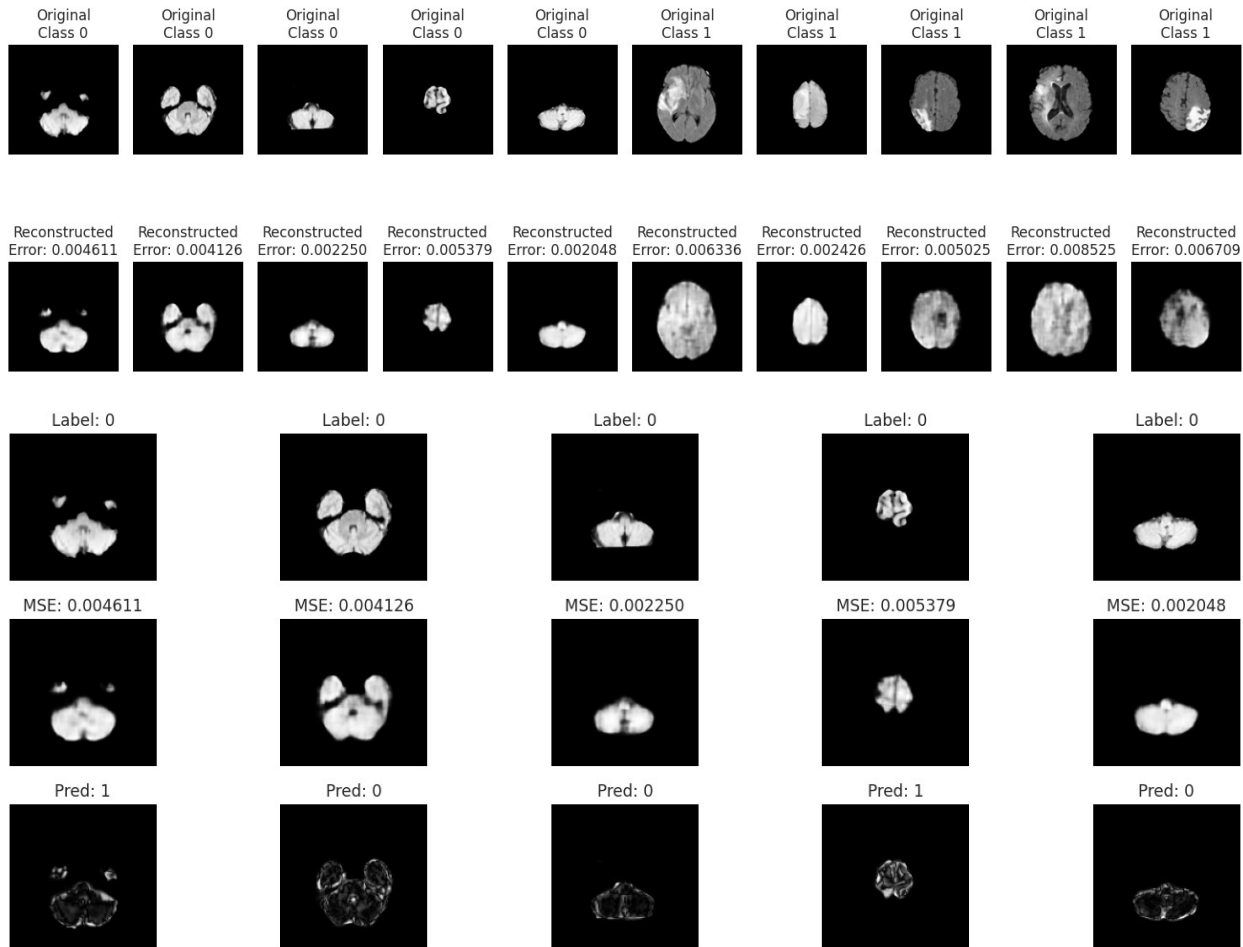
# Visualize some test images with their reconstruction and differences
(first 5 test images)
num_samples = 5
plt.figure(figsize=(15, 6))
for i in range(num_samples):
    # Original image
    plt.subplot(3, num_samples, i + 1)
    plt.imshow(test_images[i].reshape(128, 128), cmap='gray')
    plt.title(f"Label: {int(test_labels[i])}")
    plt.axis('off')

    # Reconstructed image
    plt.subplot(3, num_samples, i + 1 + num_samples)
    plt.imshow(reconstructed_images[i].reshape(128, 128), cmap='gray')
    plt.title(f"MSE: {mse[i]:.6f}")
    plt.axis('off')

    # Difference
    plt.subplot(3, num_samples, i + 1 + 2 * num_samples)
    plt.imshow(np.abs(test_images[i] -
reconstructed_images[i]).reshape(128, 128), cmap='gray')
    plt.title(f"Pred: {predictions[i]}")
    plt.axis('off')
plt.tight_layout()
plt.show()

```


36/36 _____ 0s 4ms/step
 36/36 _____ 0s 6ms/step
 47/47 _____ 0s 4ms/step
 47/47 _____ 0s 5ms/step
 Anomaly detection threshold: 0.004146
 Accuracy: 0.8172
 Confusion Matrix:
 [[480 99]
 [111 459]]



Best accuracy from VAE

```

# Reconstruct test images using VAE's encoder and decoder
_, _, z = encoder.predict(test_images) # Get latent representation
reconstructed_images = decoder.predict(z) # Decode to reconstructions

# Compute reconstruction error (MSE) for each image
mse = np.mean(np.square(test_images - reconstructed_images), axis=(1,
2, 3))

```

```

# Compute training MSE for threshold calculation
_, _, train_z = encoder.predict(train_images) # Encode training data
train_reconstructed = decoder.predict(train_z) # Decode training
reconstructions
train_mse = np.mean(np.square(train_images - train_reconstructed),
axis=(1, 2, 3))

# Dictionary to store threshold and accuracy results
results = {}

# Iterate over percentiles from 1 to 100 with tqdm
for i in tqdm(range(100), desc="Evaluating Thresholds"): # Add tqdm
with a description
    percentile = 1 + i # 1% to 100%
    threshold = np.percentile(train_mse, percentile) # Percentile of
Class 0 errors

    # Classify test images as normal (0) or anomalous (1) based on
threshold
    predictions = (mse > threshold).astype(int)

    # Evaluate performance
    accuracy = accuracy_score(test_labels, predictions)
    conf_matrix = confusion_matrix(test_labels, predictions)

    # Store results
    results[percentile] = {
        'threshold': threshold,
        'accuracy': accuracy,
        'conf_matrix': conf_matrix,
        'predictions': predictions # Store predictions for the best
result
    }

# After all iterations, print all results and find the best one
print("\n=== All Results ===")
best_percentile_vae = None
best_accuracy = -1
for percentile, result in results.items():
    if result['accuracy'] > best_accuracy:
        best_accuracy = result['accuracy']
        best_percentile_vae = percentile

print("\n=== Best Result ===")
print(f"Best Percentile: {best_percentile_vae}%")
print(f"Threshold: {results[best_percentile_vae]['threshold']:.6f}")
print(f"Accuracy: {results[best_percentile_vae]['accuracy']:.4f}")
print(f"Confusion Matrix:\n{results[best_percentile_vae]
['conf_matrix']}")

```

```
# Plot threshold vs accuracy
thresholds = [results[p]['threshold'] for p in results]
accuracies = [results[p]['accuracy'] for p in results]

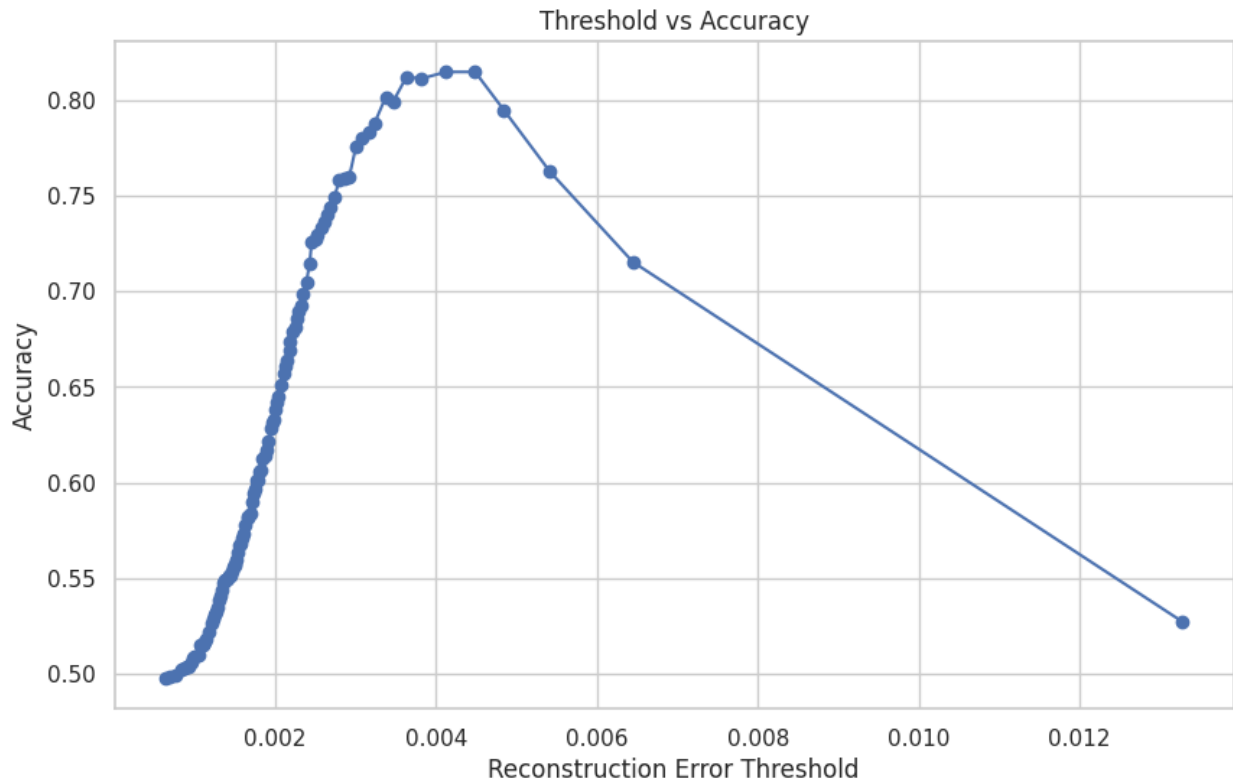
plt.figure(figsize=(10, 6))
plt.plot(thresholds, accuracies, marker='o', linestyle='--', color='b')
plt.title('Threshold vs Accuracy')
plt.xlabel('Reconstruction Error Threshold')
plt.ylabel('Accuracy')
plt.grid(True)
plt.show()
```

```
36/36 _____ 0s 3ms/step
36/36 _____ 0s 5ms/step
47/47 _____ 0s 3ms/step
47/47 _____ 0s 4ms/step
```

```
Evaluating Thresholds: 100%|██████████| 100/100 [00:00<00:00,
526.90it/s]
```

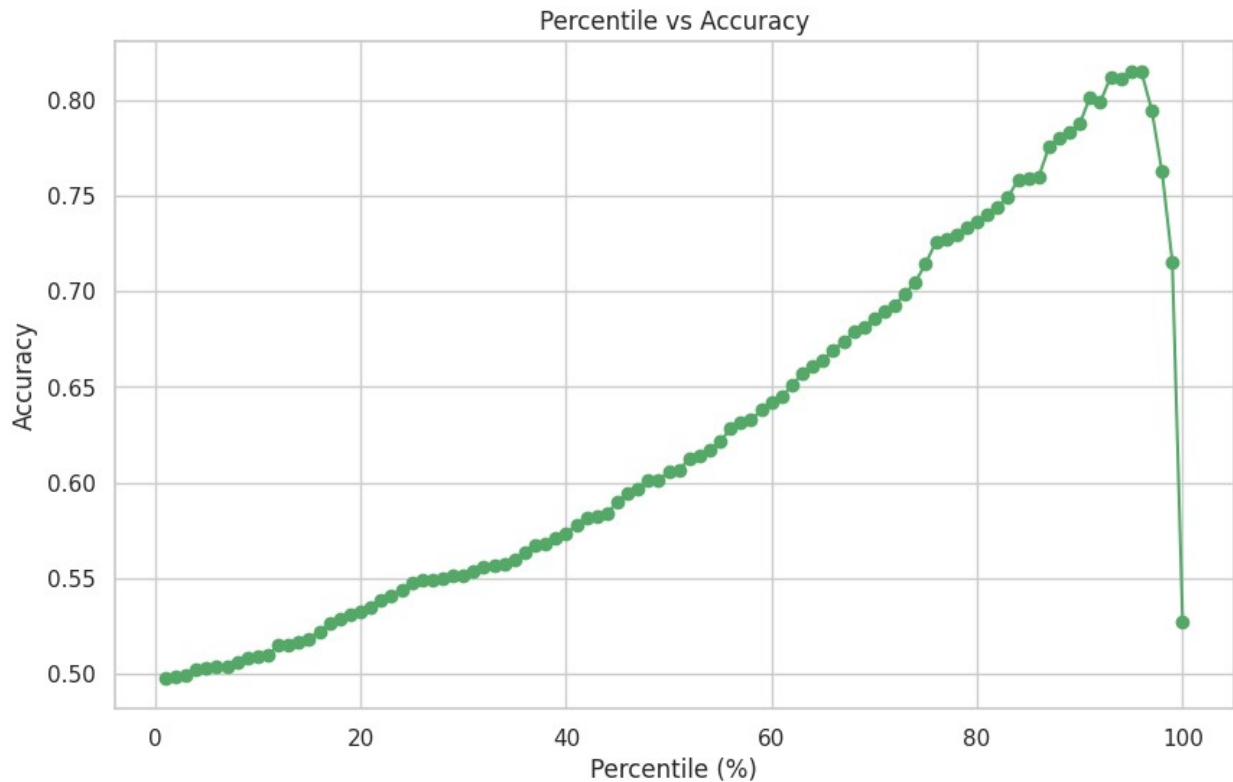
```
=== All Results ===
```

```
=== Best Result ===
Best Percentile: 95%
Threshold: 0.004125
Accuracy: 0.8146
Confusion Matrix:
[[481  98]
 [115 455]]
```



```
percentiles_vae = list(results.keys()) # [1, 2, ..., 100]
accuracies_vae = [results[p]['accuracy'] for p in results]

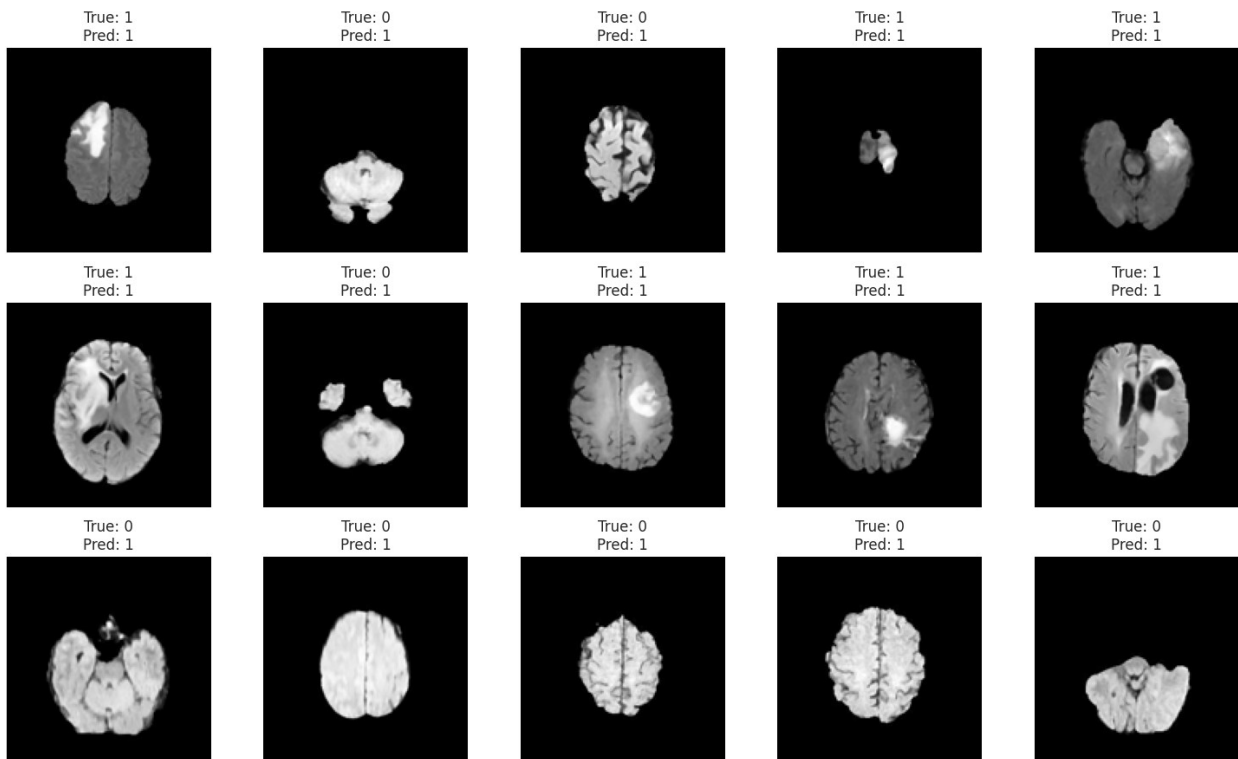
plt.figure(figsize=(10, 6))
plt.plot(percentiles_vae, accuracies_vae, marker='o', linestyle='-',
color='g')
plt.title('Percentile vs Accuracy')
plt.xlabel('Percentile (%)')
plt.ylabel('Accuracy')
plt.grid(True)
plt.show()
```



```
# Visualize predictions for the best result (15 images)
best_predictions = results[best_percentile]['predictions']
num_images_to_show = 15
indices = np.random.choice(len(test_images), num_images_to_show,
                             replace=False) # Randomly select 15 indices

plt.figure(figsize=(15, 9))
for i, idx in enumerate(indices):
    plt.subplot(3, 5, i + 1) # 3 rows, 5 columns
    plt.imshow(test_images[idx].reshape(128, 128), cmap='gray')
    plt.title(f"True: {int(test_labels[idx])}\nPred: {best_predictions[idx]}")
    plt.axis('off')
plt.tight_layout()
plt.suptitle(f"Predictions for Best Threshold (Percentile: {best_percentile_vae}%, Accuracy: {best_accuracy:.4f})", y=1.05)
plt.show()
```

Predictions for Best Threshold (Percentile: 95%, Accuracy: 0.8146)



AE vs VAE comparison

```
# Reconstruct test images using VAE's encoder and decoder
_, _, z = encoder.predict(test_images) # Get latent representation
reconstructed_images = decoder.predict(z) # Decode to reconstructions

# Compute reconstruction error (MSE) for each image
mse = np.mean(np.square(test_images - reconstructed_images), axis=(1, 2, 3))

# Compute training MSE for threshold calculation
_, _, train_z = encoder.predict(train_images)
train_reconstructed = decoder.predict(train_z)
train_mse = np.mean(np.square(train_images - train_reconstructed), axis=(1, 2, 3))

# Dictionary to store threshold and accuracy results
results_vae = {}

# Test thresholds from 1% to 100% with tqdm
for i in tqdm(range(100), desc="Analyzing"):
    percentile = 1 + i # 1% to 100%
    threshold = np.percentile(train_mse, percentile) # Percentile of Class 0 errors
```

```

    # Classify test images as normal (0) or anomalous (1) based on
    threshold
    predictions = (mse > threshold).astype(int)

    # Evaluate performance
    accuracy = accuracy_score(test_labels, predictions)
    conf_matrix = confusion_matrix(test_labels, predictions)

    # Store results
    results_vae[percentile] = {
        'threshold': threshold,
        'accuracy': accuracy,
        'conf_matrix': conf_matrix
    }

# Find the best result
best_percentile_vae = None
best_accuracy_vae = -1
for percentile, result in results_vae.items():
    if result['accuracy'] > best_accuracy_vae:
        best_accuracy_vae = result['accuracy']
        best_percentile_vae = percentile

# Plot percentile vs accuracy for VAE
percentiles_vae = list(results_vae.keys()) # [1, 2, ..., 100]
accuracies_vae = [results_vae[p]['accuracy'] for p in results_vae]
# Combined plot for comparison
plt.figure(figsize=(12, 6))
plt.plot(percentiles_ae, accuracies_ae, marker='o', linestyle='--',
         color='g', label='Autoencoder')
plt.plot(percentiles_vae, accuracies_vae, marker='o', linestyle='--',
         color='b', label='VAE')
plt.title('Percentile vs Accuracy: Autoencoder vs VAE')
plt.xlabel('Percentile (%)')
plt.ylabel('Accuracy')
plt.grid(True)
plt.legend()
plt.show()

```

```


36/36 _____ 0s 3ms/step
36/36 _____ 0s 5ms/step
47/47 _____ 0s 3ms/step
47/47 _____ 0s 4ms/step

```

```

Analyzing: 100%|██████████| 100/100 [00:00<00:00, 533.29it/s]

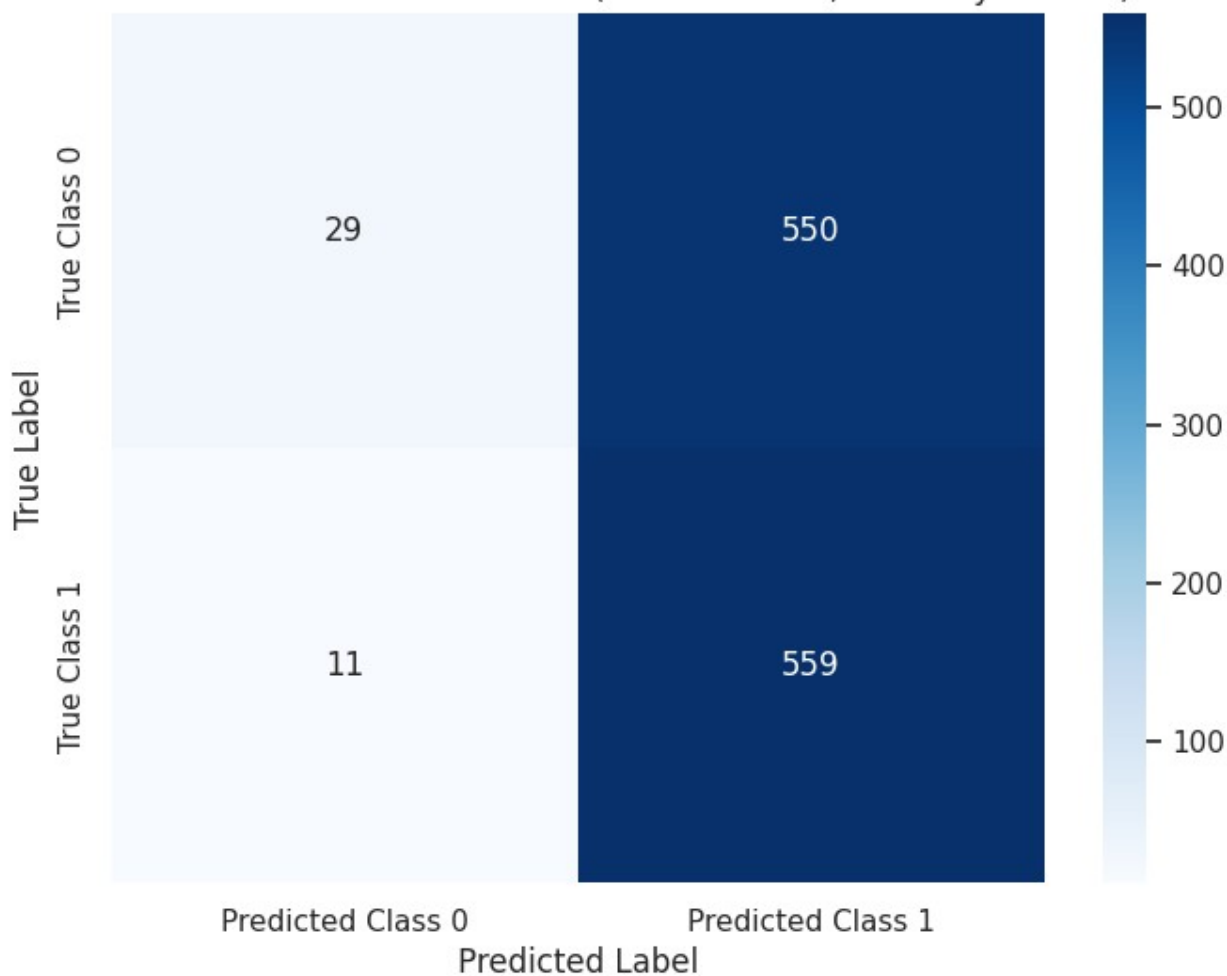
```



Heatmap comparison of best result from VAE vs AE

```
plt.figure(figsize=(8, 6))
sns.heatmap(results[best_percentile]['conf_matrix'], annot=True,
            fmt='d', cmap='Blues',
            xticklabels=['Predicted Class 0', 'Predicted Class 1'],
            yticklabels=['True Class 0', 'True Class 1'])
plt.title(f'Confusion Matrix for Best Threshold (Percentile:
{best_percentile}%, Accuracy: {best_accuracy:.4f})')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```


Confusion Matrix for Best Threshold (Percentile: 4%, Accuracy: 0.5117)



```
sns.set(style="whitegrid")
plt.figure(figsize=(8, 6))
sns.heatmap(results[best_percentile_vae]['conf_matrix'], annot=True,
            fmt='d', cmap='Blues',
            xticklabels=['Predicted Class 0', 'Predicted Class 1'],
            yticklabels=['True Class 0', 'True Class 1'])
plt.title(f'VAE Confusion Matrix (Best Percentile:
{best_percentile_vae}%, Accuracy: {best_accuracy:.4f})')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

VAE Confusion Matrix (Best Percentile: 95%, Accuracy: 0.8146)

